

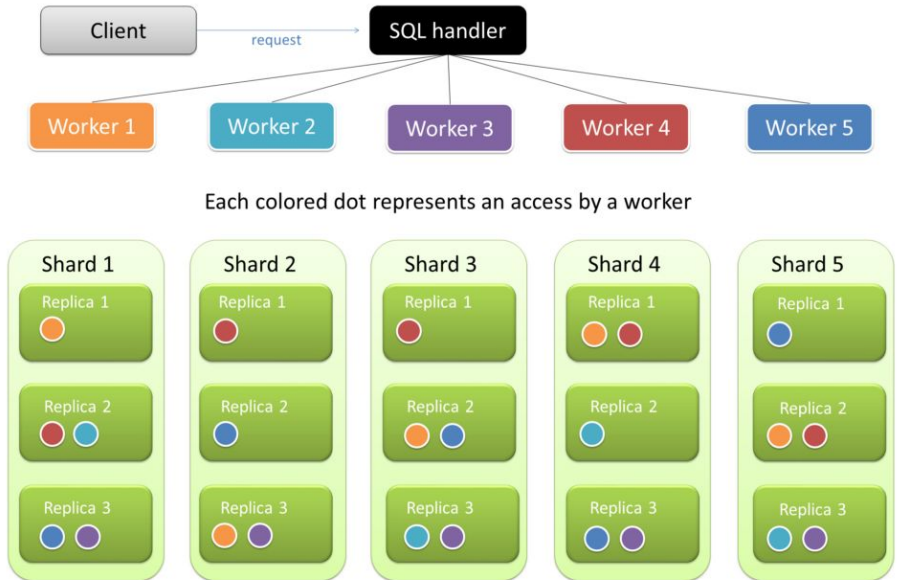



Elastic Stack

Elasticsearch, Logstash, Kibana, Beats

Elasticsearch

- Search and analytics engine (Built on Apache Lucerne)
- Document-style, schema-free (JSON) storage accessible via REST API
- Sharded horizontally (rows)
- Replicated
- Horizontally scaling
- Resilient/HA
- Distributed
- Multi-shard queries by workers (figure) - coordinated, multishard & replica querying for faster speed






All servers up and running.


P = 'Primary shard'

R = 'Replica shard'

Elasticsearch Resiliency



Server 4 goes down.

Server 2's replica of  activated as primary.

Server 6 instantiated and starts cloning data from node 5 hexagon, node 1 square, node 2 diamond.



Elasticsearch uses

- **Logs** - Import tons of log formats out-of-the-box.
- **Metrics** - Track server/docker/k8/application metrics/uptime.
- **APM - Application Performance Monitoring** - Elasticsearch help enable 'distributed tracing' which is following a client journey through different parts of the system to identify bottlenecks
- **Site, App, workplace search**
- **Maps** - visually analyzing geospatial data, identify website traffic, geospatial holes in data, outlier cases and anomalies.
- <https://www.elastic.co/elasticsearch/#what-exactly-can-i-use-elasticsearch-for%3F>



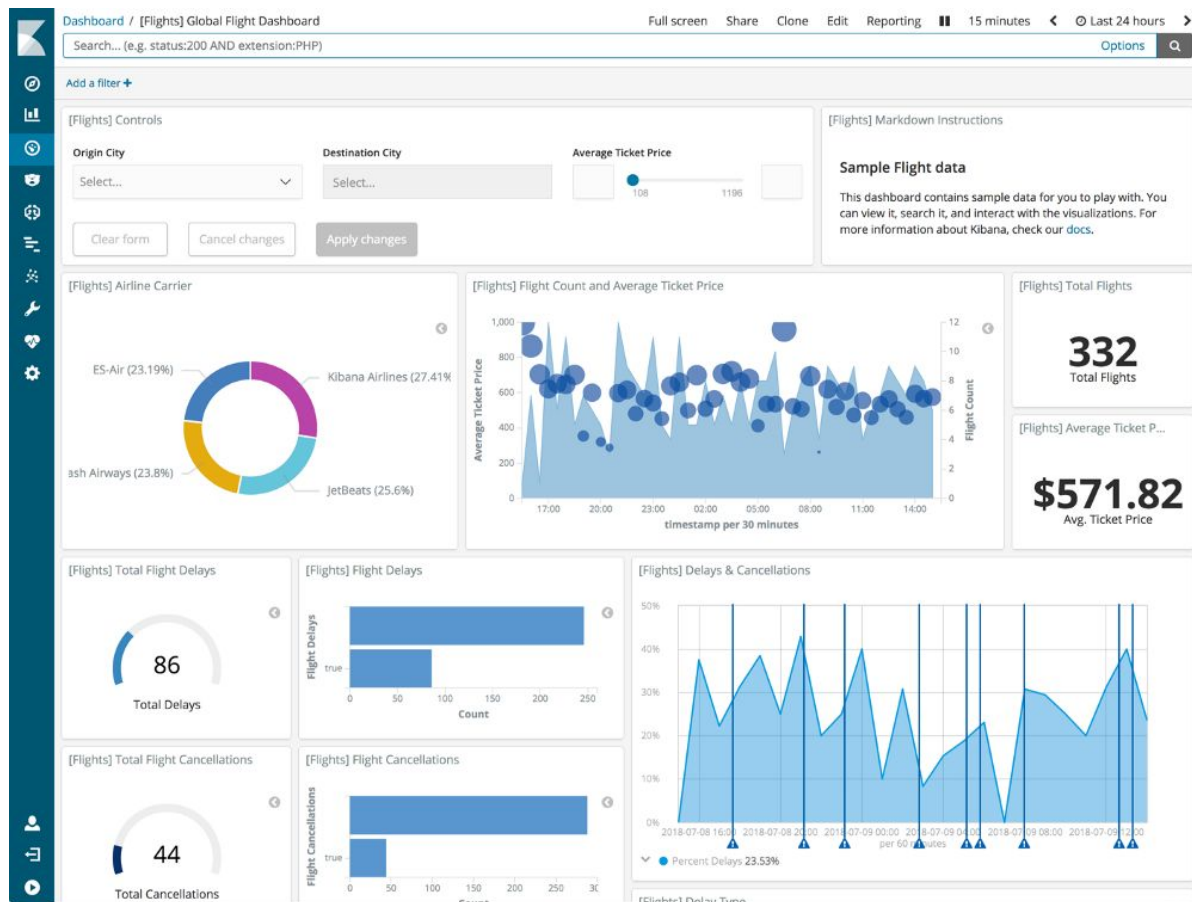
Kibana - Front-end

Provides:

1. **Admin tool** for managing an ES cluster of nodes health and performance
2. **Visualization tool** for data collected and indexed by ES

Full Feature list: <https://www.elastic.co/kibana/features>

Admin Dashboard



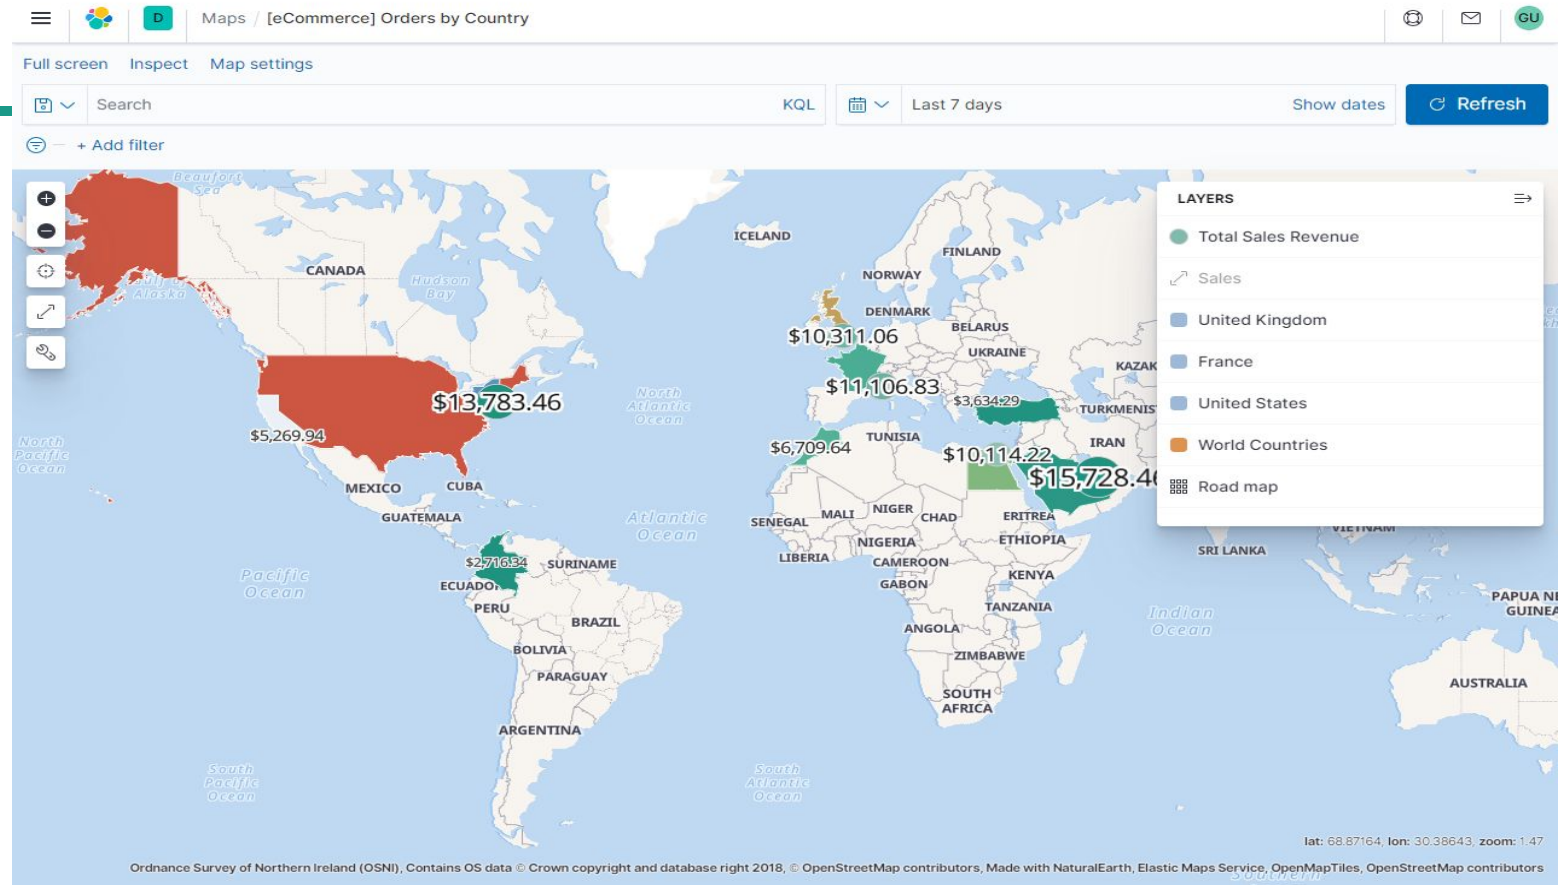
- RBAC (role-based access control)
- Alerting to email, slack, pagerDuty, webhooks
- out-of-box preconfigured for docker, dmbs, websevers
- Custom dashboards for any data



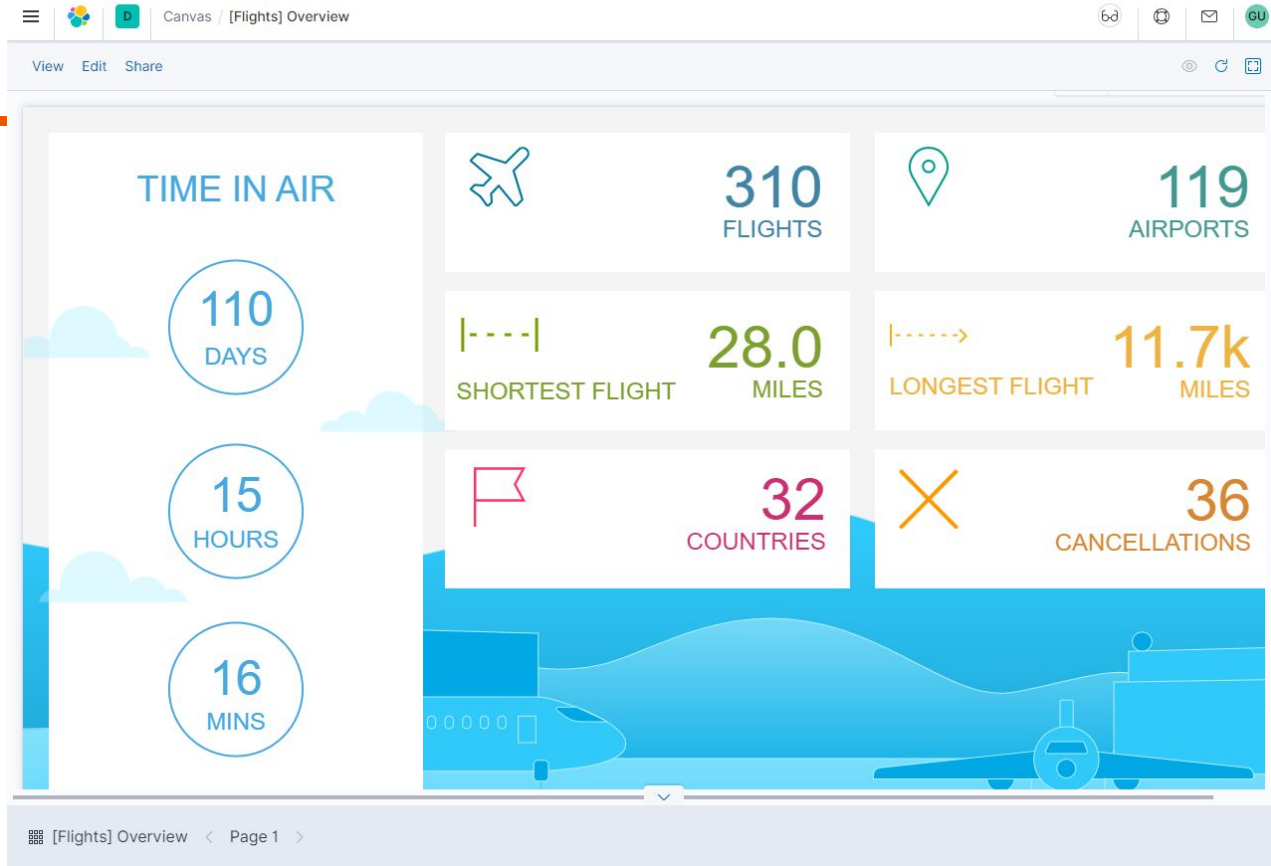
Visualizations available on Kibana

- Data can be ingested from many data sources
 - App, service, sample sets by configuring a beats data shipper
 - Upload data files (csv, json, logs) up to 100mb default (1gb max)
 - [Geospatial data](#) (lat/long columns or geoJSON/GDAL)
- [Main demo page](#) (some examples in following slides)

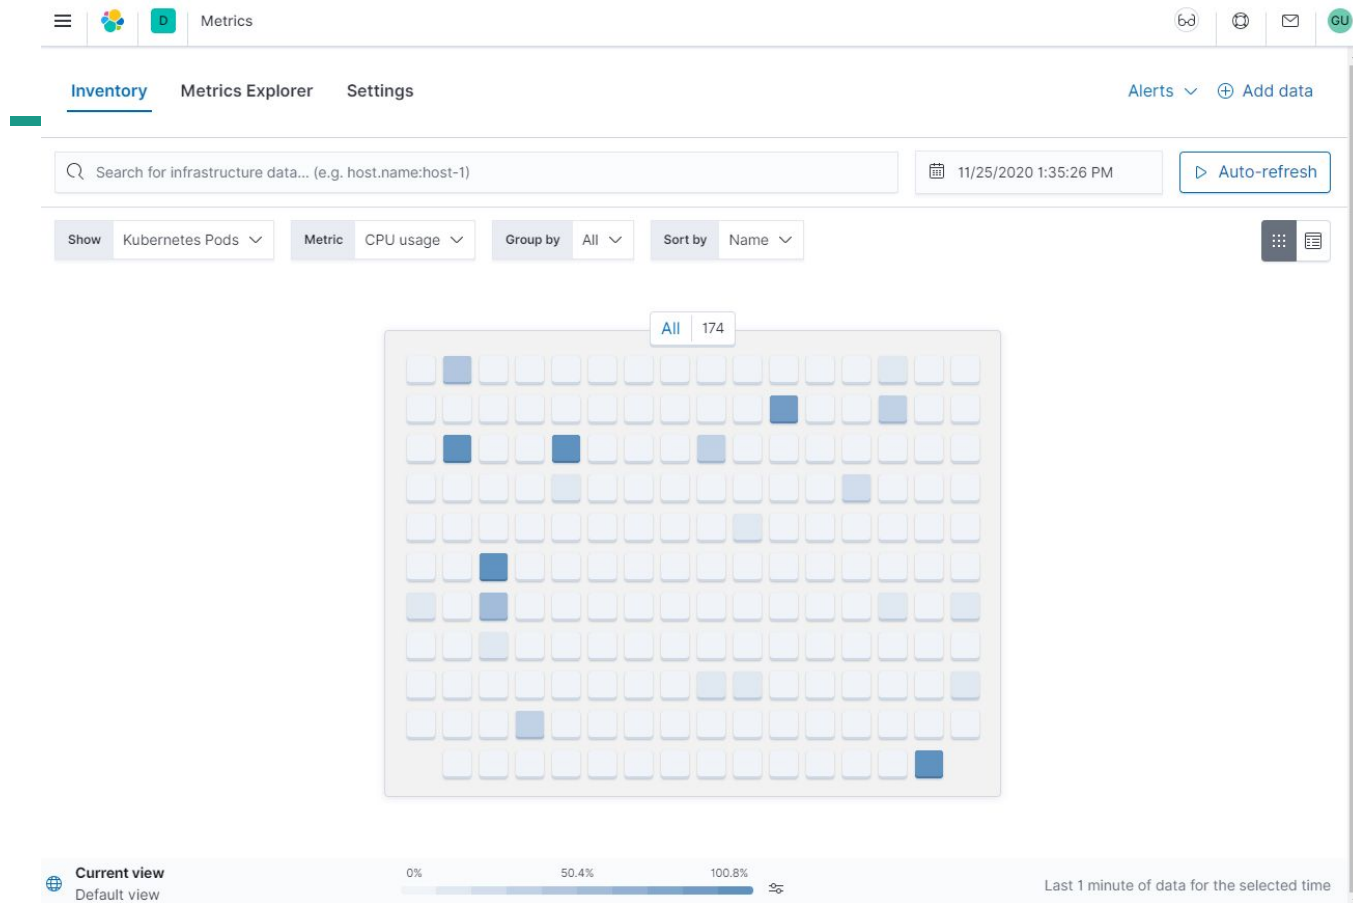
Maps demo: [\[link\]](#)



Canvas demo: [\[link\]](#) Flights



Infrastructure dashboard - k8pods - [\[link\]](#)





Beats - Data shippers

Eight beats included: (Also, [Community Beats archive for user-created Beats.](#))

- **Filebeat** - log files (backpressure handler if Logstash gets backed up). Lowers number of beats, until LS catches up.
- **Metricbeat** - Server CPU/mem/disk/network stats
- **Packetbeat** - Network packages (http), traffic, latency, error data
- **Winlogbeat** - Windows event logs, for security/activity analysis
- **Auditbeat** - Data from audit daemon, kernel level access to Linux events
- **Heartbeat** - Monitors uptime and response time for sites
- **Journalbeat** - system daemon journals for linux at all levels
- **Functionbeat** - data streaming pipeline for logs coming from FaaS (function as service aka Lambdas) platforms



Logstash

Elasticstash is the backend data store, Kibana is the front-end reporting tool. Beats can take in data and process each line, but if you want to transform that data, here is Logstash.

Logstash is where built-in inputs, filters, codecs and outputs live.

- Unstructured data => structured data
- IP addresses => geographic locations
- Fields removed, manipulated or recalculated

Many [plugins out of the box](#) for all kinds of events (beats, cloudwatch, files, gc storage, irc, kafka, redis, network stuff (tcp, snmp, udp, websocket), twitter)



Elasticsearch Modules Related to our Stack

1. Filebeat - [Kafka module](#) - The kafka module collects and parses the logs created by [Kafka](#).
2. Metricbeat - [Kafka module](#) - default (consumer, partition), also has broker, consumergroup, producer metricsets
3. Metricbeat - [ZooKeeper module](#) - default (mntr/server), also connection. Mntr = network, latency, follower, connections
4. Elasticsearch for Apache Hadoop - connector for Hadoop/hive/pig/spark/storm to Elasticsearch. [Features](#).
5. Filebeat - [nginx module](#) - parses and ingests access and error logs for nginx http server for kibana
6. Metricbeat - [nginx module](#) - periodically fetches metrics from nginx http servers (default [stubstatus](#))
7. Filebeat - [docker input](#) - reads logs from docker containers
8. Metricbeat - [docker module](#) - fetches metrics from docker containers (default: container, cpu, diskio, healthcheck, info, memory, network.)
9. Metricbeat - [k8 event metricset](#) - This is the event metricset of the Kubernetes module.
10. Metricbeat - [ec2 metricset](#) - The ec2 metricset of aws module allows you to monitor your AWS EC2 instances, including cpu, network, disk and status. ec2 metricset fetches a set of values from [Cloudwatch AWS EC2 Metrics](#).



Links


- Elastic website is [here](#). Kibana [Demo site](#)
- On AWS: [the ELK stack](#)
- On Stackshare: [Elasticsearch](#) [Kibana](#) [Logstash](#)
- For the article that a lot of this slideshow came from (somewhat rambling) check [here](#).



Kubernetes (K8), Docker, Helm



**Traditional
Deployment**




**Virtualized
Deployment**



**Container
Deployment**

**Kubernetes & Docker work
together to build & run
containerized applications**



**Kubernetes
Deployment**



Helm is a package manager and provides ready-to-deploy Instances (chart) of many popular packages. [\[link\]](#)

Chart = package

Repository = chart storage

Release = chart running in k8

Helm installs charts into Kubernetes, creating a new release for each installation. And to find new charts, you can search Helm chart repositories.



Vocab/Abbrev

- ❑ **GA** = General Availability Release (as opposed to Beta, Alpha, etc)
- ❑ **GKE** = Google Kubernetes Engine (<https://cloud.google.com/kubernetes-engine>)
- ❑ **ECK** = Elastic Cloud on Kubernetes (<https://github.com/elastic/cloud-on-k8s>)
- ❑ **Operator pattern** = <https://kubernetes.io/docs/concepts/extend-kubernetes/operator/>
- ❑ **K8 concepts**: <https://kubernetes.io/docs/concepts/>



Links

Videos

- [Introduction to Microservice, Docker and Kubernetes](#) by James Quigley(good)

References

- [Helm charts for Elastic](#) @elastic
- [All docker images](#) @elastic
- [Kubernetes icons on github](#)



Applications for PJ3



Applications for Project 3

Enterprise search

Search processed data, query same data spark is accessing by linking to HDFS/other data store that we're using as a data sink. **Logstash** ingests data ([how does it work?](#)) and outputs to **Elasticsearch**.

Observability

Monitor container, EC2, S3 metrics, logs and status ([using file, packet and metric beat](#))

Real-time visualizations

Using **Kibana**



Ways to do stuff with Logstash

[Integrations](#): pull data in through jdbc, kafka, [rabbitmq](#) (message-broker software)

[Input plugins](#): ex: beats, file, gCloud storage/pubsub, github wh, graphite, log4j, s3, twitter, http/udp, more

[Output plugins](#): csv, datadog, email, file, pubsub, mongodb.stdout/pipe, redis, s3, webhdfs, more

[Filter plugins](#): aggregate, alter, clone, csv parse, dissect unstructured data, add geo to ip data, parse json/kv/xml, mutate fields, more



Links

Slidedeck: [Building Streaming Data Pipelines with Elasticsearch, Apache Kafka and KSQL](#)

Kafka Producers->Brokers/Cluster Architecture

